

---

---

# Quantum Reinforcement Learning

Group 12

吳冠緯、吳育丞、陳宥辰

---

---


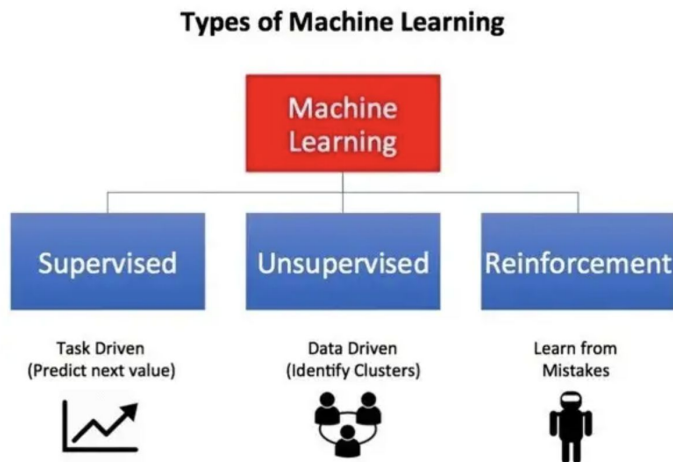
# Outline

- Value-based Reinforcement Learning
- Quantum Reinforcement Learning
- Another Quantum Application with RL

# Value-based Reinforcement Learning


# What's Reinforcement Learning?

- **Reinforcement Learning** is a method of machine learning by which an algorithm can make decisions and **take actions within a given environment**, and **learns what appropriate decisions to make through repeated trial-and-error actions**.




# Example of Reinforcement Learning

- State:
  - $s_1, s_2, \dots, s_9$
- Action:
  - Move up
  - Move down
  - Move right
  - Move left
- Reward:
  - Money: 1
  - Bomb: -1
  - Otherwise: 0



# Markov Decision Process

- A **discrete-time** stochastic control process.
- A framework for modeling decision making in situations where **the outcomes are partly random and partly under control of the decision maker**.
- **The state transitions of an MDP are independent of all previous states.**



# Modeling RL with MDP

- State (S)
- Action (A)
- Reward (R)
- Policy ( $\pi$ )

$$\pi(a | s)$$

- Objective (G)

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{T-t-1} R_T$$


# Value Function

- Formulate the objective more precisely:

$$\begin{aligned}v_{\pi}(s) &\doteq \mathbb{E}_{\pi}[G_t \mid S_t = s] \\&= \mathbb{E}_{\pi} \left[ R_{t+1} + \gamma \sum_{k=0}^{\infty} \gamma^k R_{t+k+2} \mid S_t = s \right] \\&= \mathbb{E}_{\pi} \left[ R_{t+1} + \gamma V_{\pi}(s') \mid S_{t+1} = s' \right]\end{aligned}$$

- Example:

$$V_{\pi}(S_6) = \frac{1}{4}(-1 + 0.7 * V_{\pi}(S_3)) + \frac{1}{4}(+0 + 0.7 * V_{\pi}(S_5)) + \frac{1}{4}(+0 + 0.7 * V_{\pi}(S_6)) + \frac{1}{4}(+1 + 0.7 * V_{\pi}(S_9))$$





# How to get optimal policy?

- Greedy Action

$$V_{\pi}(s) = \sum_{a \in A(s)} q_{\pi}(s, a) \Rightarrow \pi'(s) \doteq \operatorname{argmax}_a q_{\pi}(s, a)$$

$$V_{\pi}(S_6) = \downarrow \frac{1}{4}(-1 + 0.7 * V_{\pi}(S_3)) \downarrow \frac{1}{4}(+0 + 0.7 * V_{\pi}(S_5)) \downarrow \frac{1}{4}(+0 + 0.7 * V_{\pi}(S_6)) \uparrow \frac{1}{4}(+1 + 0.7 * V_{\pi}(S_9))$$

- Policy Iteration

$$\pi_0 \xrightarrow{\text{E}} v_{\pi_0} \xrightarrow{\text{I}} \pi_1 \xrightarrow{\text{E}} v_{\pi_1} \xrightarrow{\text{I}} \pi_2 \xrightarrow{\text{E}} \cdots \xrightarrow{\text{I}} \pi_* \xrightarrow{\text{E}} v_*,$$

- $\xrightarrow{\text{E}}$  稱為 policy evaluation
- $\xrightarrow{\text{I}}$  稱為 policy improvement。

→ How to accelerate the process ?

# Dynamic Programming

Input  $\pi$ , the policy to be evaluated

Algorithm parameter: a small threshold  $\theta > 0$  determining accuracy of estimation

Initialize  $V(s)$ , for all  $s \in \mathcal{S}^+$ , arbitrarily except that  $V(\text{terminal}) = 0$

Loop:

$\Delta \leftarrow 0$

Loop for each  $s \in \mathcal{S}$ :

$v \leftarrow V(s)$

$V(s) \leftarrow \sum_a \pi(a|s) \sum_{s', r} p(s', r | s, a) [r + \gamma V(s')]$

$\Delta \leftarrow \max(\Delta, |v - V(s)|)$

until  $\Delta < \theta$

→ Can we always know the value function in advance ?

# Monte Carlo Method and Temporal Difference Learning

$$V_{n+1}(s) = \mathbb{E}_{\pi}[G_t \mid S_t = s]$$
$$= \frac{G_{t1} + G_{t2} + G_{t3} + \dots + G_{tn}}{n}$$

$$= \frac{1}{n}(G_{tn} + (n-1)\frac{1}{n-1} \sum_{i=1}^{n-1} G_{ti})$$

$$= \frac{1}{n}(G_{tn} + (n-1)V(s))$$

$$= \frac{1}{n}(G_{tn} + nV(s) - V(s))$$

$$= V_n(s) + \frac{1}{n} [G_{tn} - V_n(s)]$$

$$V(S_t) \leftarrow V(S_t) + \alpha [G_t - V(S_t)]$$

$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

Final formulation of updating value function

# Quantum Reinforcement Learning

# Define Actions with Qubit States

- An action set can be represented by a superposition state of  $n$  qubits.

$$N_a \leq 2^n \leq 2N_a \quad |A\rangle = \sum_n \beta_n |a_n\rangle \quad \sum_n |\beta_n|^2 = 1$$

- When we measure the qubit, it will collapse into one of its eigen action  $|a_n\rangle$  with the probability of  $|\beta_n|^2$ .

# Exploitation vs Exploration in RL


- Greedy action (exploitation)

$$\pi'(s) \doteq \operatorname{argmax}_a q_{\pi}(s, a)$$

- Epsilon-greedy (exploration)

A probability of the agent randomly conducting an action rather than greedy action.

→ But how?



# Utilize Grover's Algorithm

- Start with equally weighted superposition (n Hadamard gates)

$$|a_0^{(n)}\rangle = \frac{1}{\sqrt{2^n}} \left( \sum_{a=00\dots 0}^{\overbrace{11\dots 1}^n} |a\rangle \right)$$

- Amplitude amplification

$$\pi'(s) \doteq \operatorname{argmax}_a q_\pi(s, a)$$

$$\Rightarrow U_a = I - 2|a\rangle\langle a|$$

$$\Rightarrow U_{a_0^{(n)}} = H^{\otimes n}(2|0\rangle\langle 0| - I)H^{\otimes n} = 2|a_0^{(n)}\rangle\langle a_0^{(n)}| - I$$

# QRL Algorithm

## Procedural QRL:

Initialize  $|s^{(m)}\rangle = \sum_{s=00\dots 0}^{11\dots 1} C_s |s\rangle$ ,  $f(s) = |a_s^{(n)}\rangle = \sum_{a=00\dots 0}^{11\dots 1} C_a |a\rangle$  and  $V(s)$  arbitrarily

Repeat (for each episode)

For all states  $|s\rangle$  in  $|s^{(m)}\rangle = \sum_{s=00\dots 0}^{11\dots 1} C_s |s\rangle$ :

1. Observe  $f(s) = |a_s^{(n)}\rangle$  and get  $|a\rangle$ ;
2. Take action  $|a\rangle$ , observe next state  $|s'\rangle$ , reward  $r$ , then

(a) Update state value:  $V(s) \leftarrow V(s) + \alpha(r + \gamma V(s') - V(s))$

(b) Update probability amplitudes:

repeat  $U_{Grover}$  for  $L$  times

$$U_{Grover} |a_s^{(n)}\rangle = U_{a_0^{(n)}} U_a |a_s^{(n)}\rangle$$

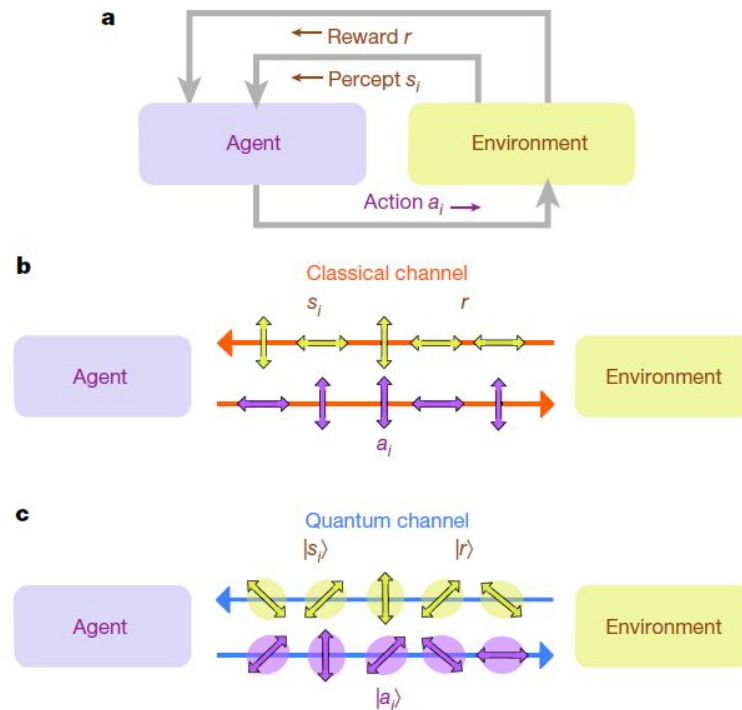
Until for all states  $|\Delta V(s)| \leq \varepsilon$ .



# Another Quantum Application with RL

# Schematic of Learning Agent

- Agent and environment interacting **classically**, where communication is only possible via a fixed preferred basis.
- Agent and environment interacting via a **quantum** channel, where arbitrary **superposition states** are exchanged.

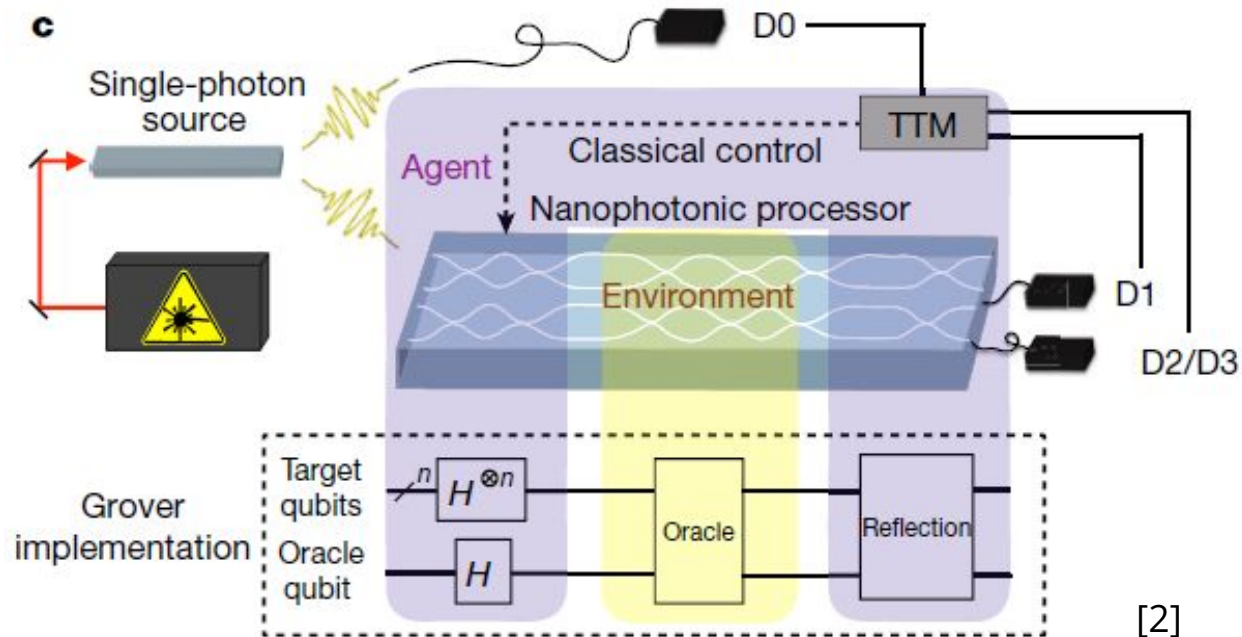


# Framework

- Focus on so-called **deterministic strictly epochal** (DSE) learning scenarios.
- Here 'epochs' consist of strings of percepts  $\mathbf{s} = (\mathbf{s}_0, \dots, \mathbf{s}_{L-1})$  with fixed  $s_0$ , actions  $\mathbf{a} = (\mathbf{a}_1, \dots, \mathbf{a}_L)$  of fixed length  $L$ , and a final reward  $\mathbf{r}$ , and both  $\mathbf{s} = \mathbf{s}(\mathbf{a})$  and  $\mathbf{r} = \mathbf{r}(\mathbf{a})$  are completely determined by  $\mathbf{a}$ .


$$U_E |\mathbf{a}\rangle_A |0\rangle_R = \begin{cases} |\mathbf{a}\rangle_A |1\rangle_R & \text{if } r(\mathbf{a}) > 0 \\ |\mathbf{a}\rangle_A |0\rangle_R & \text{if } r(\mathbf{a}) = 0 \end{cases}.$$

# Experiment Setup



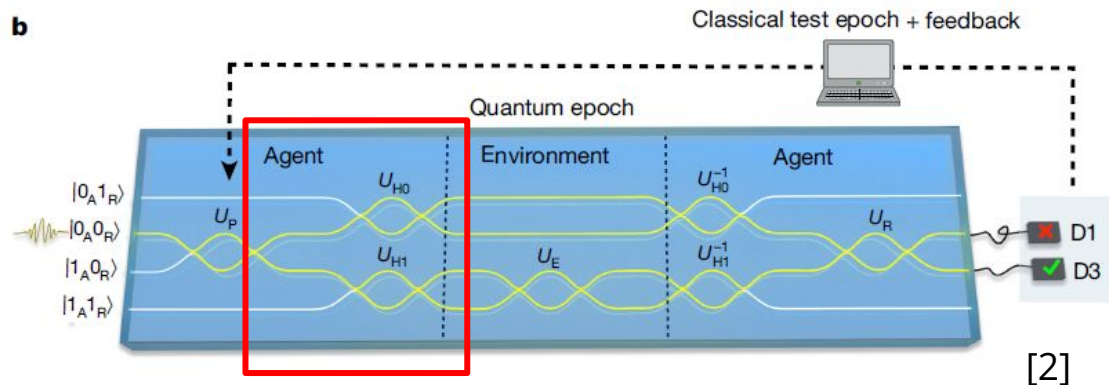
# Classical Epoch

- In a classical strategy, the environment **flips** the reward qubit only if the action qubit is in the winning state via  $U_E$ . Next, the photon is **coupled out and detected** in either D1 or D2 with probability  $\cos^2(\xi)$  and  $\sin^2(\xi)$ , respectively.



# Quantum Epoch

(1) The agent prepares the state  $|\psi\rangle_A |-\rangle_R$ , with  $|\psi\rangle_A = \sum_a \sqrt{p(\mathbf{a})} |\mathbf{a}\rangle_A = \cos(\xi)|\ell\rangle_A + \sin(\xi)|w\rangle_A$ , and sends it to the environment.  $|w\rangle_A$  and  $|\ell\rangle_A$  are superpositions of all winning (rewarded) and losing (non-rewarded) action sequences, respectively, and  $|-\rangle_R = (|0\rangle_R - |1\rangle_R)/\sqrt{2}$ .




# Quantum Epoch

(2) The environment applies  $U_E$  from equation (1) to  $|\psi\rangle_A |-\rangle_R$ , flipping the sign of the winning state:


$$U_E |\psi\rangle_A |-\rangle_R = [\cos(\xi)|\ell\rangle_A - \sin(\xi)|w\rangle_A] |-\rangle_R, \quad (3)$$

and returns the resulting state to the agent.



# Quantum Epoch


(3) The agent performs a reflection  $U_R = 2|\psi\rangle\langle\psi|_A - \mathbb{1}_A$  over the initial state  $|\psi\rangle_A$ .






# Improvement in Learning Time

- For  $\varepsilon < P$ ,  $\eta_Q > \eta_C$ , meaning that the quantum strategy proves advantageous over the classical case.
- However, as soon as  $\eta_Q = \eta_C$  (at  $P = 0.396$ ), a classical agent starts outperforming a quantum-enhanced agent that still performs Grover iterations.



[2]

# Combined Strategy



# Topics Candidates

- Implement QRL algorithm in work 1 on different RL tasks.
- Combine the advantages of previous two works on certain RL tasks.
- Implement QNN for RL tasks.

# Reference

- [1] Dong, D.; Chen, C.; Li, H.; Tarn, T.-J. Quantum Reinforcement Learning. IEEE Trans. Syst. Man Cybern. Part B (Cybern.) 2008, 38, 1207.
- [2] Saggio, V., Asenbeck, B.E., Hamann, A. *et al.* Experimental quantum speed-up in reinforcement learning agents. *Nature* 591, 229–233 (2021).